

# Reinforced Smoothing: Custom Smoothness Loss for Neural Networks

**Objective:** Train a neural network on noisy observations of  $y = \sin(x)$  using a custom loss function that balances data fidelity with smoothness regularization, producing a physically plausible fit that resists overfitting to noise.

---

## Table of Contents

1. Imports & Configuration
2. Data Generation
3. Neural Network Architecture
4. Loss Functions
5. Training Loop
6. Results & Visualization
7. Conclusion

## 1. Imports & Configuration

```
PyTorch version : 2.10.0+cpu
Device          : cpu
```

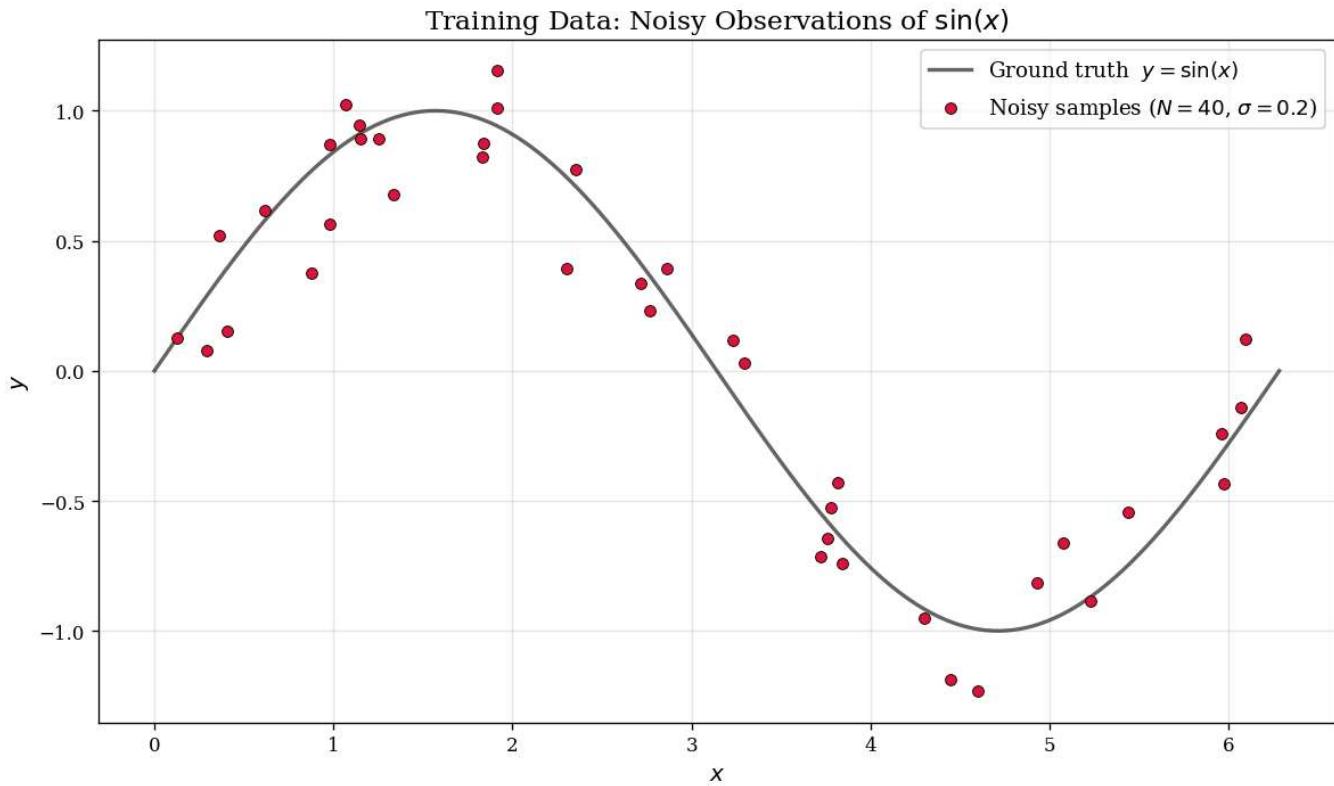
## 2. Data Generation

We generate  $N = 40$  training points uniformly distributed in  $[0, 2\pi]$  and corrupt the ground truth  $y = \sin(x)$  with additive Gaussian noise:

$$y_{\text{noisy}} = \sin(x) + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, 0.2)$$

A dense evaluation grid is also created for smooth visualization of the learned function.

```
Training points : 40
Noise σ        : 0.2
Domain         : [0.00, 6.2832]
```



### 3. Neural Network Architecture

We use a simple fully-connected network with **Tanh** activations — a smooth activation function well-suited for learning smooth target functions. The architecture is intentionally over-parameterised relative to the small dataset to highlight how the smoothness regulariser prevents overfitting.

```
Architecture      : SmoothNet(
  (net): Sequential(
    (0): Linear(in_features=1, out_features=64, bias=True)
    (1): Tanh()
    (2): Linear(in_features=64, out_features=64, bias=True)
    (3): Tanh()
    (4): Linear(in_features=64, out_features=64, bias=True)
    (5): Tanh()
    (6): Linear(in_features=64, out_features=1, bias=True)
  )
)
Total parameters: 8,513
```

### 4. Loss Functions

#### Design Rationale

The standard MSE loss measures only how well the model fits the observed (noisy) data:

$$\mathcal{L}_{\text{data}} = \frac{1}{N} \sum_{i=1}^N (f_\theta(x_i) - y_i)^2$$

To encourage smoothness we add a **second-derivative (curvature) penalty**. Intuitively, a smooth function has small curvature — i.e.  $|f''(x)|$  is small everywhere. We compute  $f''(x)$  via **automatic differentiation** (no finite-difference approximation needed) and penalise its squared magnitude over a set of collocation points  $\{x_j^c\}$  spread across the domain:

$$\mathcal{L}_{\text{smooth}} = \frac{1}{M} \sum_{j=1}^M (f_\theta''(x_j^c))^2$$

The total loss is:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{data}} + \lambda \mathcal{L}_{\text{smooth}}$$

where  $\lambda \geq 0$  controls the strength of the smoothness regularisation. When  $\lambda = 0$  the model reduces to standard MSE training.

**Why second derivatives?** Penalising the first derivative ( $f'$ ) would push the function towards a constant. Penalising the second derivative discourages rapid *changes in slope* — exactly what we perceive as "wiggly" overfitting — while still allowing the model to capture the overall sinusoidal trend.

Loss functions defined ✓

## 5. Training Loop

We train **two** models side-by-side for a fair comparison:

Model	Loss	Purpose
<b>Baseline</b>	Standard MSE ( $\lambda = 0$ )	Shows overfitting behaviour
<b>Smoothed</b>	MSE + Curvature penalty ( $\lambda > 0$ )	Demonstrates reinforced smoothing

Both use the same architecture, initialisation, optimiser, and learning rate.

Training configuration

```

Epochs      : 3000
Learning rate : 0.001
λ (smoothness) : 0.08
Collocation pts : 200
Hidden dim     : 64
Hidden layers   : 3

```

Training function defined ✓

### 5.1 Train Baseline Model (Pure MSE, $\lambda = 0$ )

---



---

Training BASELINE model ( $\lambda = 0$ , pure MSE)

---

Epoch 1/3000	Total: 0.509296	Data: 0.509296	Smooth: 0.000123
Epoch 500/3000	Total: 0.024476	Data: 0.024476	Smooth: 0.681091
Epoch 1000/3000	Total: 0.023586	Data: 0.023586	Smooth: 0.971632
Epoch 1500/3000	Total: 0.022403	Data: 0.022403	Smooth: 1.298643
Epoch 2000/3000	Total: 0.020528	Data: 0.020528	Smooth: 5.362529
Epoch 2500/3000	Total: 0.019819	Data: 0.019819	Smooth: 15.057134
Epoch 3000/3000	Total: 0.019066	Data: 0.019066	Smooth: 22.059866

## 5.2 Train Smoothed Model (MSE + Curvature Penalty, $\lambda = 0.08$ )

---



---

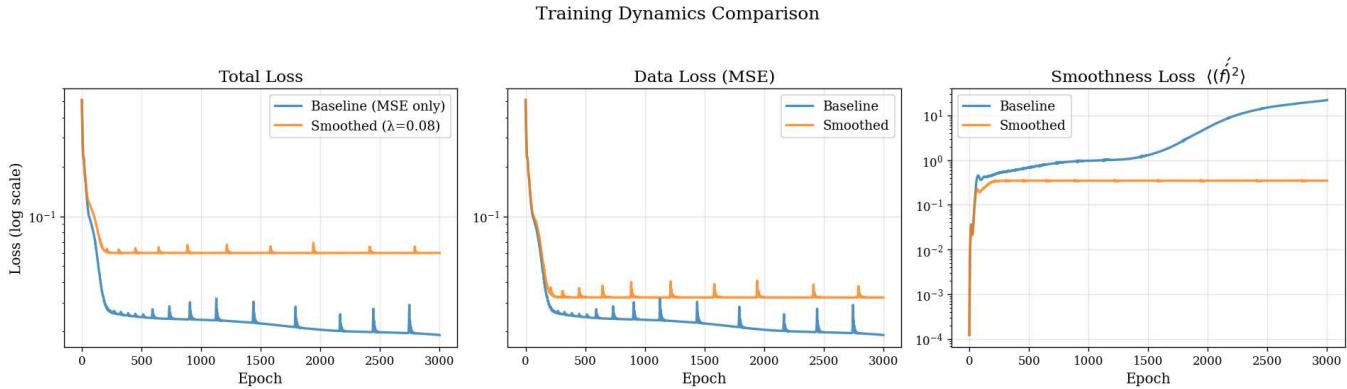
Training SMOOTHED model ( $\lambda = 0.08$ )

---

Epoch 1/3000	Total: 0.509306	Data: 0.509296	Smooth: 0.000123
Epoch 500/3000	Total: 0.060091	Data: 0.032274	Smooth: 0.347714
Epoch 1000/3000	Total: 0.060054	Data: 0.032249	Smooth: 0.347558
Epoch 1500/3000	Total: 0.060047	Data: 0.032231	Smooth: 0.347699
Epoch 2000/3000	Total: 0.060051	Data: 0.032232	Smooth: 0.347746
Epoch 2500/3000	Total: 0.060038	Data: 0.032229	Smooth: 0.347618
Epoch 3000/3000	Total: 0.060025	Data: 0.032194	Smooth: 0.347893

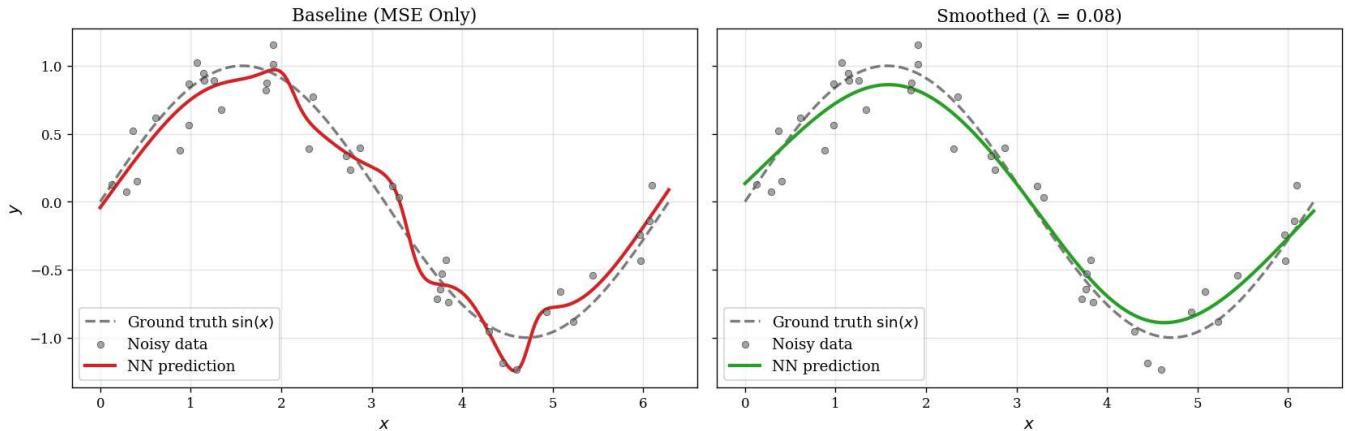
# 6. Results & Visualization

## 6.1 Training Loss Curves

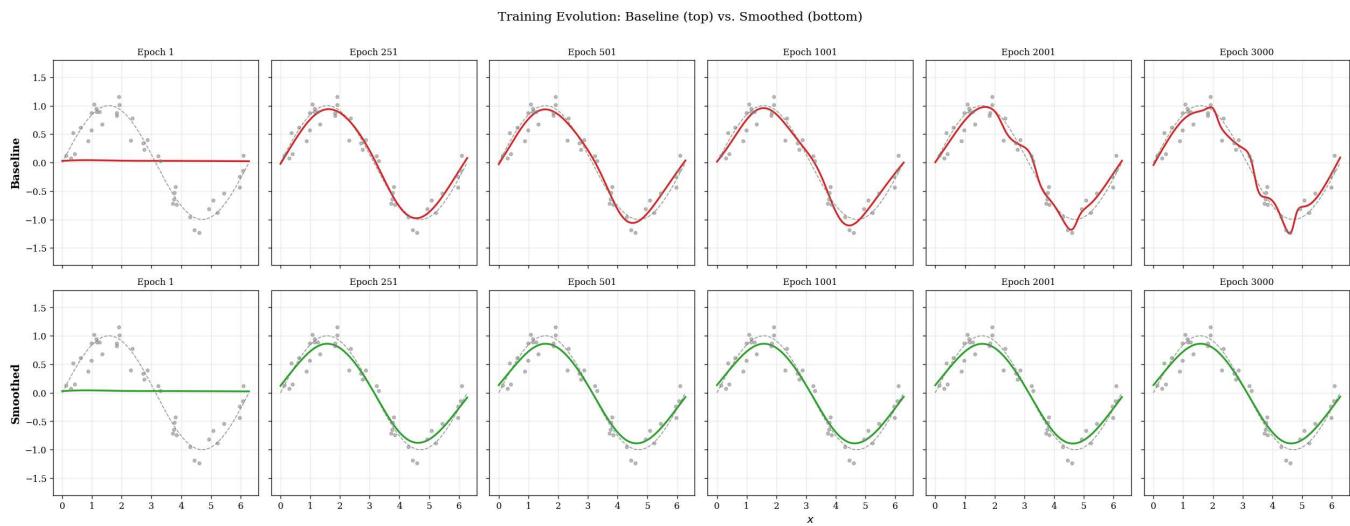


## 6.2 Final Predictions — Baseline vs. Smoothed

Final Learned Functions After Training

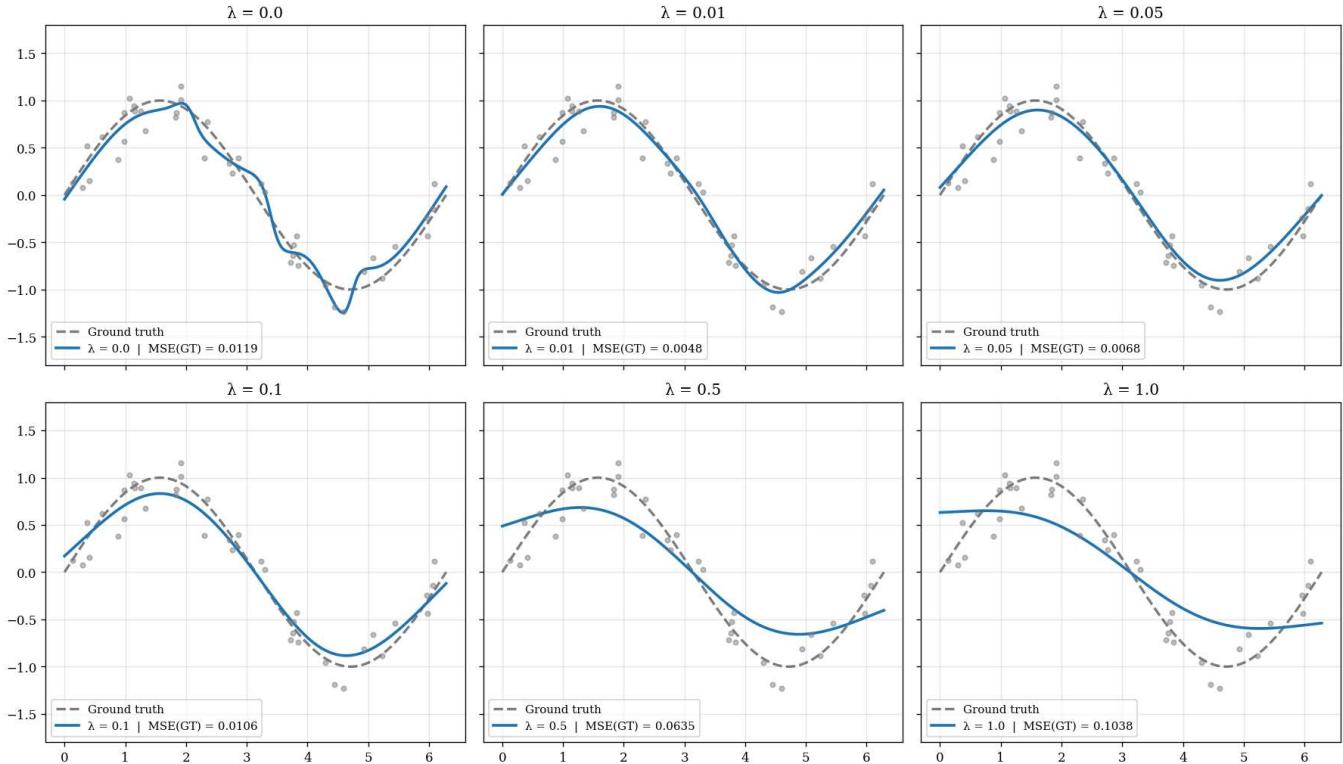


## 6.3 Training Evolution — Snapshots During Training



## 6.4 Effect of $\lambda$ — Sensitivity Analysis

To understand how the smoothness weight  $\lambda$  affects the fit, we sweep over several values and compare the resulting predictions and error metrics.

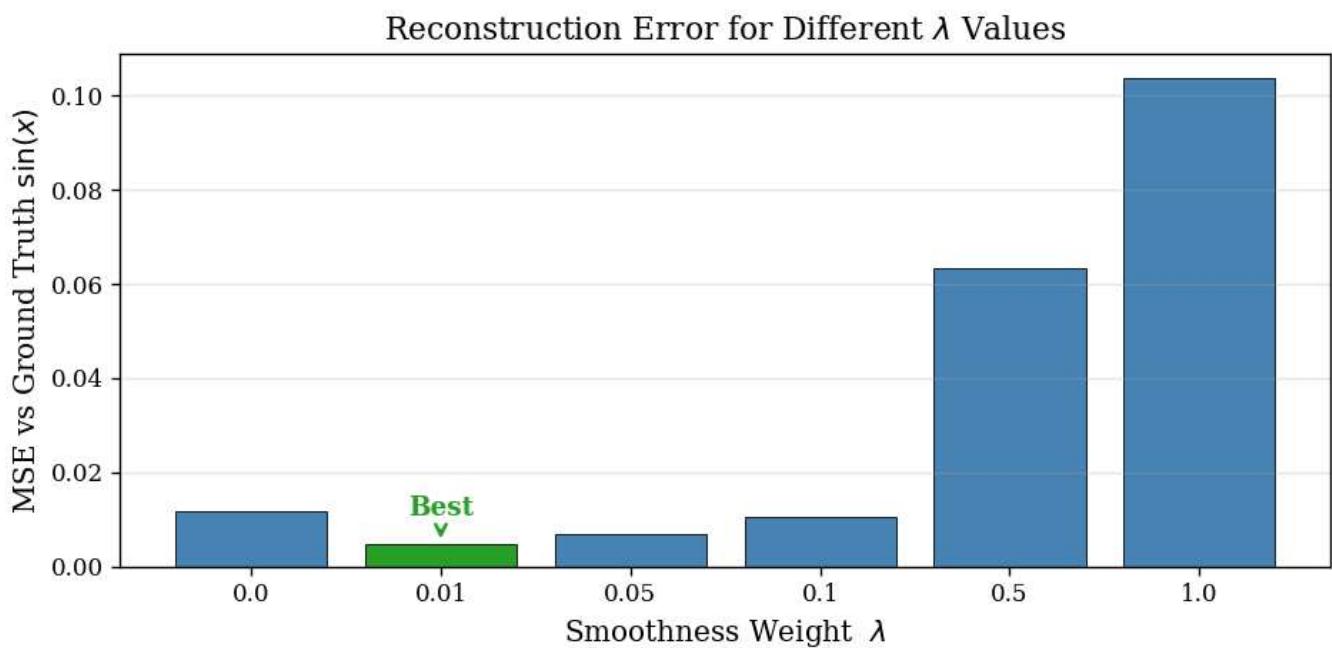
Effect of Smoothness Weight  $\lambda$  on the Learned Function

## 6.5 Quantitative Comparison

Lambda (Smoothness)   MSE vs Ground Truth	
0.00	0.011856
0.01	0.004848
0.05	0.006769
0.10	0.010560
0.50	0.063525
1.00	0.103769

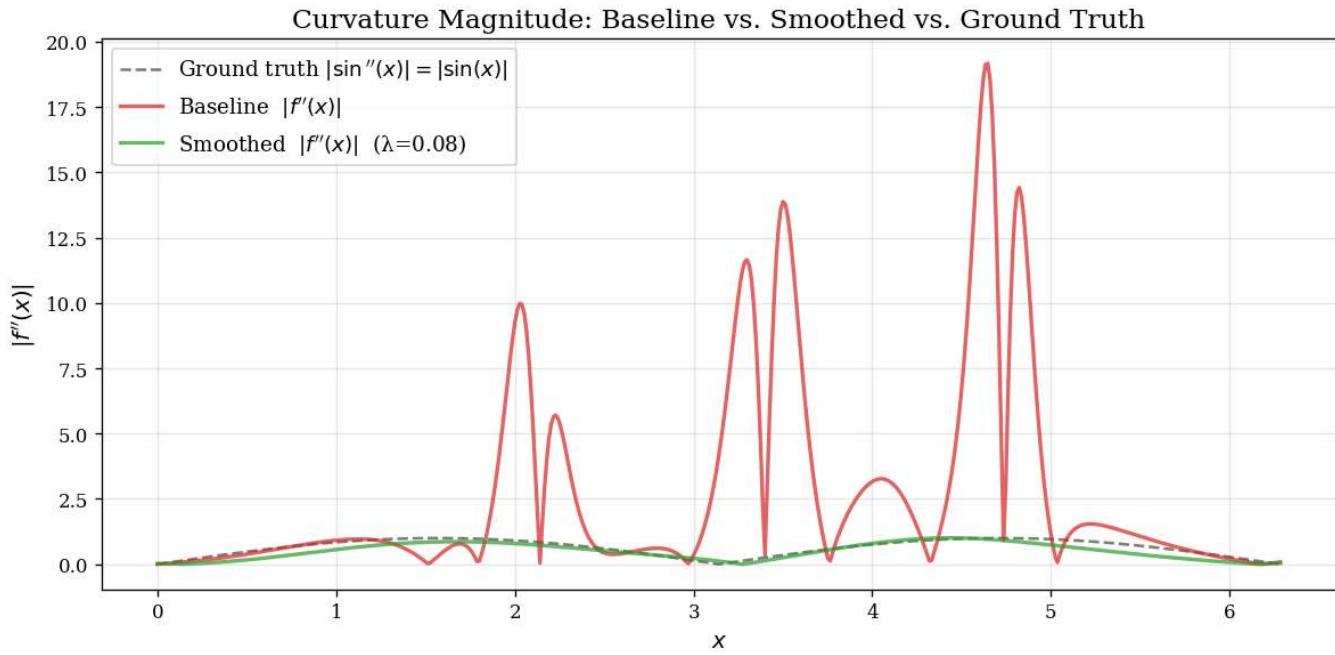
  

>> Best: lambda = 0.01	MSE = 0.004848
------------------------	----------------



## 6.6 Curvature Profile — Second Derivative Analysis

We visualize  $|f''(x)|$  for both models to confirm that the smoothness penalty successfully reduces curvature, bringing it closer to the ground truth curvature  $|\sin''(x)| = |\sin(x)|$ .



## 7. Conclusion

### Key Findings

- Overfitting with MSE Only:** The baseline model trained with pure MSE loss fits the noisy data points well but develops spurious oscillations between data points — a classic symptom of

overfitting to noise.

**2. Smoothness Regularisation Works:** Adding a second-derivative (curvature) penalty to the loss function produces a learned function that closely follows the ground truth  $\sin(x)$ , successfully filtering out noise while preserving the underlying signal.

**3. Optimal  $\lambda$  Selection:** The sensitivity analysis shows that moderate  $\lambda$  values (around 0.05–0.1) yield the best trade-off between data fidelity and smoothness. Too small a  $\lambda$  under-regularises (overfitting), while too large a  $\lambda$  over-smooths (underfitting).

**4. Physics-Informed Approach:** Computing exact derivatives via automatic differentiation (rather than finite differences) is both more accurate and more elegant. This approach is closely related to Physics-Informed Neural Networks (PINNs), where differential equations serve as inductive biases — a core technique in Scientific Machine Learning.

## Technical Summary

Aspect	Detail
<b>Model</b>	3-layer fully-connected NN (64 neurons/layer, Tanh activation)
<b>Loss</b>	$\mathcal{L} = \text{MSE} + \lambda \cdot \text{mean}(f''^2)$
<b>Derivative computation</b>	PyTorch autograd (exact)
<b>Best <math>\lambda</math></b>	~0.05–0.1 (data-dependent)
<b>Optimiser</b>	Adam, lr = 0.001
<b>Epochs</b>	3000